

Synchronization

(OpenMP Synchronization)

Synchronization

- OMP_SINGLE
- OMP MASTER
- OMP CRITICAL
- OMP ATOMIC
- OMP BARRIER
- OpenMP programs use shared variables to communicate.
- We need to make sure these variables are not accessed at the same time by different threads (will cause race conditions, WHY?).
- OpenMP provides a number of directives for synchronization.

OMP_SINGLE

- Only one of the threads will execute the following block of code
 - **The rest will wait for it to complete**
 - Good for non-thread-safe regions of code (such as I/O)
 - Must be used in a parallel region
 - **Applicable to parallel for sections**

OMP_SINGLE

- It specifies that the enclosed code is to be **executed by only one thread in the team.**
- The thread chosen could vary from one run to another.
- Threads that are not executing in the SINGLE directive wait at the END SINGLE directive unless NOWAIT is specified.

```
#include<stdio.h>
#include<omp.h>
int main()
{ int i, sal1=0,sal2=0;
#pragma omp parallel shared(sal1,sal2)
{
for(i=0;i<5;i++)
{ Sal1=sal1+i; }
Printf("Salary of the first company=%d", sal1);
for(i=0;i<5;i++)
{
Sal2=sal2+i
}
Printf("Salary of the first company=%d", sal2);
}
Return 0;
}
```

Threads – 4

Data Races

```
Salary First Company = 10      thread 1
Salary second Company = 10     thread 1
Salary First Company = 30      thread 3
Salary second Company = 20     thread 3
Salary First Company = 40      thread 2
Salary second Company = 30     thread 2
Salary First Company = 20      thread 0
Salary second Company = 40     thread 0
admin@ubuntu ~ $
```

```
#include<stdio.h>
#include<omp.h>
Int main()
{ Int l, sal1=0,sal2=0;
#pragma omp parallel reduction(:+sal1)
For(i=0;i<5;i++)
{
Sal1=sal1+i
}
Printf("Salary of the first company=%d", sal1);
#pragma omp parallel reduction(:+sal2)
For(i=0;i<5;i++)
{
Sal2=sal2+i
}
Printf("Salary of the first company=%d", sal2);
Return 0;
}
```

```
Salary First Company = 10          thread no 0  
Salary second Company = 10        thread no 0  
admin@ubuntu ~ $
```

- This solution is quite nice. **It doesn't have any data races** and it uses less lines of code to express the solution.
- A possible disadvantage of this version is that we **have to wait for computations of all salaries to complete before we are able to see the results.**


```
#pragma omp parallel shared(sal1,sal2)
{
#pragma omp parallel reduction(:+sal1)
For(i=0;i<5;i++)
{Sal1=sal1+i}
#pragma omp single
{ Printf("Salary of the first company=%d", sal1); }
#pragma omp parallel reduction(:+sal2)
For(i=0;i<5;i++)
{
Sal2=sal2+i
}}
Printf("Salary of the first company=%d", sal2);
Return 0; }
```

```
admin@ubuntu:~$ ./8.001
```

```
Salary First Company = 10      thread no 1
```

```
Salary second Company = 10     thread no 1
```

```
Salary second Company = 20     thread no 0
```

```
Salary second Company = 30     thread no 3
```

```
Salary second Company = 40     thread no 2
```

```
#pragma omp parallel shared(sal1,sal2)
{
#pragma omp parallel reduction(:+sal1)
For(i=0;i<5;i++)
{Sal1=sal1+i}
#pragma omp single
{
    Printf("Salary of the first company=%d", sal1);
}
#pragma omp parallel reduction(:+sal2)
For(i=0;i<5;i++)
{Sal2=sal2+i }
#pragma omp single
{
    Printf("Salary of the first company=%d", sal2);
} }
Return 0;
}
```

OMP_SINGLE

Only one
thread
initializes the
shared
variable a

```
#pragma omp parallel shared(a,b) private(i)
{
    #pragma omp single
    {
        a = 10;
        printf("Single construct executed by thread %d\n",
            omp_get_thread_num());
    }
    /* A barrier is automatically inserted here */

    #pragma omp for
    for (i=0; i<n; i++)
        b[i] = a;

} /*-- End of parallel region --*/

printf("After the parallel region:\n");
for (i=0; i<n; i++)
    printf("b[%d] = %d\n",i,b[i]);
```

OMP Master

- **Specifies a region** that is **to be executed only by the master thread of the team.**
- **All other threads on the team skip this section** of code
- It is similar to the SINGLE construct.
- This Directive ensures that only the master threads executes instructions in the block.
- **There is no implicit barrier so other threads will not wait for master to finish.**
- The following block will be executed by the master thread
- **No synchronization involved**
- Applicable only to parallel sections

OMP Master

```
#include <omp.h>
#include <stdio.h>
int main( )
{   int a[5], i;
#pragma omp parallel
{ // Perform some computation.
    #pragma omp for
    for (i = 0; i < 5; i++)
a[i] = i * i;
    // Print intermediate results.
    #pragma omp master
    for (i = 0; i < 5; i++)
    printf("a[%d] = %d\n", i, a[i]);
    // Wait.
    #pragma omp barrier // Continue with the computation.
    #pragma omp for
    for (i = 0; i < 5; i++) a[i] += i; } }
```

```
admin@ubuntu ~ $ ./a.out
```

```
a[3]= 9   Thread 2
```

```
a[4]= 16           Thread 2
```

```
a[0]= 0   Thread 2
```

```
a[1]= 1   Thread 2
```

```
a[2]= 4   Thread 2
```

```
a[0]=0    thread 1
```

```
a[1]=1    thread 1
```

```
a[2]=4    thread 1
```

```
a[3]=9    thread 1
```

```
a[4]=16   thread 1
```

```
a[0]=0    thread 1
```

```
a[1]=2    thread 1
```

```
a[2]=6    thread 1
```

```
a[3]=12   thread 1
```

```
a[4]=20   thread 1
```

```
admin@ubuntu ~ $
```

Difference

Single
Master

1.Wait

1.No wait

2.Separate thread for that region

2.Master thread

OMP CRITICAL

- **This Directive makes sure that only one thread can execute the code in the block.**
- **If another threads reaches the critical section it will wait untill the current thread finishes this critical section.**
- It provides a means to ensure that multiple threads do not attempt to update the same shared data simultaneously.
- An optional name can be given to a critical construct. Name must be global and unique
- When a thread encounters a critical construct, it waits until no other thread is executing a critical region with the same name.
- Every thread will execute the critical block and they will synchronize at end of critical section

```

#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
#define SIZE 10
int main()
{ int i; int max; int a[SIZE];
  for (i = 0; i < SIZE; i++)
  { a[i] = rand();
    printf(" a[%d] = %d \t thread no %d\n", i,
a[i]),omp_get_num_threads()); }
  max = a[0];
  #pragma omp parallel for
  for (i = 1; i < SIZE; i++) {
    if (a[i] > max)
    { #pragma omp critical
    { // compare a[i] and max again because max could have been
changed by another thread after
// the comparison outside the critical section
if (a[i] > max) max = a[i]; }
    } }
    printf("max = %d\t thread no = %d\n", max,
omp_get_num_threads()); }

```

Output

41 18467 6334 26500 19169 15724 11478

20258 26062 24464 max = 20258

OMP CRITICAL

#pragma omp critical [name]

- Standard critical section functionality
- **Critical sections are global in the program**
 - Can be used to protect a single resource in different functions
- Critical sections are identified by the name
 - **All the unnamed critical sections are mutually exclusive throughout the program**
 - All the critical sections having the same name are mutually exclusive between themselves

```
i n t x = 0 ;
```

```
#pragma omp p a r a l l e l s h a r e d ( x )
```

```
{ #pragma omp c r i t i c a l
```

Difference between single and critical

- single specifies that a section of code should be executed **by single thread**(not necessarily the master thread)
- critical specifies that code is executed by one thread at a time

Example

```
int a=0, b=0;
#pragma omp parallel num_threads(4)
{
    #pragma omp single
    a++;
    #pragma omp critical
    b++;
}
printf("single: %d -- critical: %d\n", a, b);
```

```
admin@ubuntu ~$ gcc -fopenmp -DNP_SINGLE
admin@ubuntu ~$ ./a.out
single 1          ---Critical 4
admin@ubuntu ~$
```

OMP ATOMIC

- `#OMP ATOMIC`
- This Directive is very similar to the `#OMP CRITICAL`
- Difference is that `#OMP ATOMIC` is only used for the update of a memory location.
- Specifies that a specific memory location must be updated atomically, rather than letting multiple threads attempt to write to it.
- Read,Write,Update,Capture

OMP ATOMIC

- The **omp atomic** directive allows access of a specific memory location atomically.
- It ensures that race conditions are avoided through direct control of concurrent threads that might read or write to or from the particular memory location.

```
for(i = 0; i<10; i++)  
{  
    #pragma omp atomic  
    x[j[i]] = x[j[i]] + 1.0;  
}
```


Ordered

#pragma omp ordered

- It allows one to **execute a structured block within a parallel loop in sequential order.**
- The **code outside this block runs in parallel.**
- if threads finish out of order, there may be an additional performance penalty because some threads might have to wait.

Ordered

```
#pragma omp parallel for  
for(i = 0; i < nproc; i++)  
{  
    do_lots_of_work(result[i]);
```

```
    #pragma omp ordered  
    fprintf(fid, "%d %f\  
n, "i, result[i]);  
    #pragma omp end  
ordered
```

OMP BARRIER

Important restrictions:

- Each barrier must be encountered by all threads in a team, or by none at all.
- The sequence of work-sharing regions and barrier regions encountered must be the same for every thread in the team.
- Without these rules some threads wait forever (or until somebody kills the process) for other threads to reach a barrier

OMP BARRIER

- **Synchronizes all threads in the team.**
- When a BARRIER directive is reached, **a thread will wait at that point until all other threads have reached that barrier.**
- All threads then resume executing in parallel the code that follows the barrier.

OMP BARRIER

- **barrier** synchronizes threads

```
#pragma omp parallel
private(myid,istart,iend)
myrange(myid,istart,iend);
for(i=istart; i<=iend; i++){
    a[i] = a[i] - b[i];
}
#pragma omp barrier
myval[myid] = a[istart] + a[0]
```

- Here **barrier** assures that $a(1)$ or $a[0]$ is available before computing **myval**

OMP BARRIER

- Performs a barrier synchronization between all the threads in a team at the given point.
- !\$OMP BARRIER will **enforce every thread to wait at the barrier until all threads have reached the barrier.**
- !\$OMP BARRIER is probably the most well known synchronization mechanism; **explicitly or implicitly.**
- The following omp directives we discussed before include an implicit barrier:

REDUCTION

```
f o r ( j =0; j<N; j++) {  
  sum = sum+a [ j ] b [ j ] ;  
}
```

- How to parallelize this code?
 - sum is not private, but accessing it atomically is too expensive
 - Have a private copy of sum in each thread, then add them up
- Use the reduction clause!

#pragma omp parallel for reduction(+: sum)

- Any associative operator must be used: +, -, ||, |, *, etc.
- The private value is initialized automatically (to 0, 1, ~0 . . .)

Exercise

- Develop an OpenMP C Parallel Program to perform the travelling salesman problem