

20BCE1550
Samridh Anand Paatni
CSE4001 Lab 10
MPI Collective Communication

a:
Code:

```
ex10a.c > main(int, char * [])
1  /*
2  a. Write a program in MPI to generate 'n' random float
3  ... numbers and send 'k' of those to each node and make
4  ... them compute the average and send it back to the
5  ... master which computes the average of those averages.
6  */
7  #include <stdio.h>
8  #include <stdlib.h>
9  #include <mpi.h>
10
11 int main(int argc, char *argv[]) {
12     int procRank, procNum;
13     int n, k, numSend, numRecv;
14     double avg;
15     MPI_Init(&argc, &argv);
16     MPI_Comm_size(MPI_COMM_WORLD, &procNum);
17     MPI_Comm_rank(MPI_COMM_WORLD, &procRank);
18
19     k = 4;
20     n = k * procNum;
21
22     if (procRank == 0) {
23         for (int dest = 1; dest < procNum; dest++) {
24             for (int j = 0; j < k; j++) {
25                 numSend = rand() % 100;
26                 printf("send: P%d(%d) -> P%d\n", procRank,
27                        numSend, dest);
28                 MPI_Send(&numSend, 1, MPI_INT, dest, 1,
29                        MPI_COMM_WORLD);
29             }
29         }
29     }
```

```

30     ...} else {
31         ... int count = 0, sum = 0, from = 0;
32         ... for (int i = 0; i < k; i++) {
33             ... MPI_Recv(&numRecv, 1, MPI_INT, from, 1,
34                 ... MPI_COMM_WORLD, MPI_STATUS_IGNORE);
35             ... count++;
36             ... sum += numRecv;
37             ... printf("recv: P0->P%d(%d) [%d/%d]\n", procRank,
38                 ... numRecv, count, k);
39         ... }
40     ...
41     ... MPI_Barrier(MPI_COMM_WORLD);
42     ...
43     ... if (procRank != 0) {
44         ... MPI_Send(&avg, 1, MPI_DOUBLE, 0, 1, MPI_COMM_WORLD);
45         ... printf("P%d(%f) ->P0\n", procRank, avg);
46     ... } else {
47         ... int count = 0;
48         ... double sum = 0.0;
49         ... for (int i = 1; i < procNum; i++) {
50             ... MPI_Recv(&avg, 1, MPI_DOUBLE, i, 1,
51                 ... MPI_COMM_WORLD, MPI_STATUS_IGNORE);
52             ... printf("P%d->P0(%f) [%d/%d]\n", i, avg, count,
53                 ... procNum-1);
54             ... count++;
55             ... sum += avg;

```

```

56         ... }
57         ... printf("\nFinal average = %f\n", sum / count);
58     ... }
59     ... MPI_Finalize();
60     ... return 0;
61 }

```

Output:

```
mpicc ex10a.c -o ex10a.out
mpiexec -np 4 ex10a.out
send: P0(83)->P1
send: P0(86)->P1
send: P0(77)->P1
send: P0(15)->P1
send: P0(93)->P2
send: P0(35)->P2
send: P0(86)->P2
send: P0(92)->P2
send: P0(49)->P3
send: P0(21)->P3
send: P0(62)->P3
send: P0(27)->P3
recv: P0->P1(83) [1/4]
recv: P0->P1(86) [2/4]
recv: P0->P1(77) [3/4]
recv: P0->P1(15) [4/4]
calc: P1_average=65.250000
recv: P0->P2(93) [1/4]
recv: P0->P2(35) [2/4]
recv: P0->P2(86) [3/4]
recv: P0->P2(92) [4/4]
calc: P2_average=76.500000
recv: P0->P3(49) [1/4]
recv: P0->P3(21) [2/4]
recv: P0->P3(62) [3/4]
recv: P0->P3(27) [4/4]
calc: P3_average=39.750000
P3(39.750000)->P0
P2(76.500000)->P0
P1->P0(65.250000) [0/3]
P2->P0(76.500000) [1/3]
P3->P0(39.750000) [2/3]

Final average = 60.500000
P1(65.250000)->P0
```


b:

Code:

```
ex10b.c > main(int, char * [])
1  /*
2   b. Write a MPI program to compute PI using "dartboard"
3   technique for 1000 rounds by using reduction
4   collective computation.
5   */
6
7  #include <stdio.h>
8  #include <stdlib.h>
9  #include <time.h>
10 #include <mpi.h>
11
12 double getRandom() {
13     return (double)(rand() % 1000000) / 500000 - 1;
14 }
15
16 int main(int argc, char *argv[]) {
17     srand(clock());
18     int procRank, procNum;
19     MPI_Init(&argc, &argv);
20     MPI_Comm_size(MPI_COMM_WORLD, &procNum);
21     MPI_Comm_rank(MPI_COMM_WORLD, &procRank);
22
23     int count = 0, total = 0;
24     int finalCount = 0, finalTotal = 0;
25
26     if (procRank != 0) {
27         for (int i = 0; i < 10000; i++) {
28             double x = getRandom();
29             double y = getRandom();
30             if ((x*x) + (y*y) < 1) count++;
31             total++;
```

```

30     .... if ((x*x)+(y*y)<1) count++;
31     .... total++;
32     .... }
33     .... printf("P%d: %d/%d\n", procRank, count, total);
34     .... }
35     .... MPI_Reduce(&count, &finalCount, 1, MPI_INT, MPI_SUM, 0,
36     .... MPI_COMM_WORLD);
37     .... MPI_Reduce(&total, &finalTotal, 1, MPI_INT, MPI_SUM, 0,
38     .... MPI_COMM_WORLD);
39     .... if (procRank == 0) {
40     ....     printf(
41     ....         "Throws inside circle = %d\nTotal throws = %d\npi
42     ....         = %f\n",
43     ....         finalCount,
44     ....         finalTotal,
45     ....         (double)finalCount/finalTotal * 4
46     ....     );
47     .... }
48     .... MPI_Finalize();
49 }

```

Output:

```

[~/vit/CSE4001_Parallel-and_Distributed_Systems]$
[~/vit/CSE4001_Parallel-and_Distributed_Systems]$ (20:00:00) —> make ex10b
mpicc ex10b.c -o ex10b.out
mpiexec -np 4 ex10b.out
P2: 7837/10000
P3: 7886/10000
P1: 7895/10000
Throws inside circle = 23618
Total throws = 30000
pi = 3.149067
[~/vit/CSE4001_Parallel-and_Distributed_Systems]$

```

c:

Code:

/*

c. Write a MPI program to perform matrix multiplication
(1000x1000) using scatter and gather routines.

*/

```
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>
#include <mpi.h>
#include <stdio.h>
```

```
#define SIZE 8
```

```
int A[SIZE][SIZE], B[SIZE][SIZE], C[SIZE][SIZE];
```

```
void fill_matrix(int m[SIZE][SIZE])
{
    static int n=0;
    int i, j;
    for (i=0; i<SIZE; i++)
        for (j=0; j<SIZE; j++)
            m[i][j] = n++;
}
```

```
void print_matrix(int m[SIZE][SIZE])
{
    int i, j = 0;
    for (i=0; i<SIZE; i++) {
        printf("\n\t");
        for (j=0; j<SIZE; j++)
            printf("%2d ", m[i][j]);
        printf("\n");
    }
}
```

```
int main(int argc, char *argv[])
{
    int myrank, P, from, to, i, j, k;
    int tag = 666; /* any value will do */
    MPI_Status status;
    MPI_Init (&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank); /* who am i */
    MPI_Comm_size(MPI_COMM_WORLD, &P); /* number of processors */

    /* Just to use the simple variants of MPI_Gather and MPI_Scatter we */
    /* impose that SIZE is divisible by P. By using the vector versions, */
    /* (MPI_Gatherv and MPI_Scatterv) it is easy to drop this restriction. */
}
```

```

if (SIZE%P!=0) {
if (myrank==0) printf("Matrix size not divisible by number of processors\n");
MPI_Finalize();
exit(-1);
}

```

```

from = myrank * SIZE/P;
to = (myrank+1) * SIZE/P;

```

```

/* Process 0 fills the input matrices and broadcasts them to the rest */
/* (actually, only the relevant stripe of A is sent to each process) */

```

```

if (myrank==0) {
fill_matrix(A);
fill_matrix(B);
}

```

```

MPI_Bcast (B, SIZE*SIZE, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Scatter (A, SIZE*SIZE/P, MPI_INT, A[from], SIZE*SIZE/P, MPI_INT, 0, MPI_COMM_WORLD);

```

```

printf("computing slice %d (from row %d to %d)\n", myrank, from, to-1);
for (i=from; i<to; i++)
for (j=0; j<SIZE; j++) {
C[i][j]=0;
for (k=0; k<SIZE; k++)
C[i][j] += A[i][k]*B[k][j];
}

```

```

MPI_Gather (C[from], SIZE*SIZE/P, MPI_INT, C, SIZE*SIZE/P, MPI_INT, 0, MPI_COMM_WORLD);

```

```

if (myrank==0) {
printf("\n\n");
print_matrix(A);
printf("\n\n\t * \n");
print_matrix(B);
printf("\n\n\t = \n");
print_matrix(C);
printf("\n\n");
}

```

```

MPI_Finalize();
return 0;
}

```

Output:

```
~/vit/CSE4001_Parallel-and-Distributed-Computing_ETLP/Lab/MPI-
(20:00:09)→ make ex10c
mpicc ex10c.c -o ex10c.out
mpiexec -np 4 ex10c.out
computing slice 0 (from row 0 to 1)

      |  0  1  2  3  4  5  6  7 |
      |  8  9 10 11 12 13 14 15 |
      | 16 17 18 19 20 21 22 23 |
      | 24 25 26 27 28 29 30 31 |
      | 32 33 34 35 36 37 38 39 |
computing slice 1 (from row 2 to 3)
computing slice 2 (from row 4 to 5)
computing slice 3 (from row 6 to 7)
      | 40 41 42 43 44 45 46 47 |
      | 48 49 50 51 52 53 54 55 |
      | 56 57 58 59 60 61 62 63 |

      *

      | 64 65 66 67 68 69 70 71 |
      | 72 73 74 75 76 77 78 79 |
      | 80 81 82 83 84 85 86 87 |
      | 88 89 90 91 92 93 94 95 |
      | 96 97 98 99 100 101 102 103 |
      | 104 105 106 107 108 109 110 111 |
      | 112 113 114 115 116 117 118 119 |
      | 120 121 122 123 124 125 126 127 |

      =

      | 2912 2940 2968 2996 3024 3052 3080 3108 |
      | 8800 8892 8984 9076 9168 9260 9352 9444 |
      | 14688 14844 15000 15156 15312 15468 15624 15780 |
      | 20576 20796 21016 21236 21456 21676 21896 22116 |
      | 26464 26748 27032 27316 27600 27884 28168 28452 |
      | 32352 32700 33048 33396 33744 34092 34440 34788 |
      | 38240 38652 39064 39476 39888 40300 40712 41124 |
      | 44128 44604 45080 45556 46032 46508 46984 47460 |

~/vit/CSE4001_Parallel-and-Distributed-Computing_ETLP/Lab/MPI-
(20:14:29)→
```